Artificial Intelligence

Lecture 2, Chapter 3
Solving Problems by Searching

Supta Richard Philip

Department of CS
AIUB

AIUB, January 2024

Table of Contents

Table of Contents

Problem Solving Agent

- A problem-solving agent is one kind of goal-based agent.
- Lets consider, the route-finding problem in the cities of Romania.
- This problem can be solved by searching.

Problem formulation

- Intelligent agents are supposed to act in such a way that the environment goes through a sequence of states that **maximizes the performance measure**.

- We will consider **a goal to be a set of states** - just those states in which the goal is satisfied.

- **Actions** can be viewed as causing **transitions between states**.

- Problem formulation is the process of deciding what actions and states to consider.

Table of Contents

Search Problem formulation

A search problem consists of:

- A state space
- A successor function(with action, cost)
- A start state and a goal test
- A solution is a sequence of actions (a plan) which transforms the start state to a goal state

Example- Vacuum World

- In this case there are eight possible world states.
- There are three possible actions: left, right, and suck.
- The goal is to clean up all the dirt, i.e., the goal is equivalent to the set of states 7,8.

Example - Route-Finding in Romania

- initial states: Arad
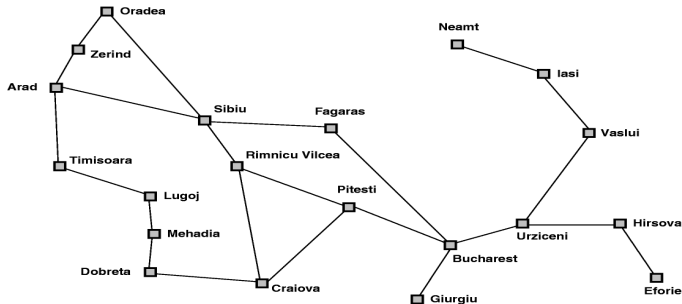- goal state: Bucharest
- operators: successor function $S(x)$ set of possible actions
- path cost: a function that assigns a cost to a path.

Example - The 8-puzzle problem

5	4	
6	1	8
7	3	2

Start State

| 1	2	3
8		4
7	6	5

Goal State

- states: a state description specifies the location of each of the eight tiles in one of the nine squares. For efficiency it is also useful to include a location for the blank.
- operators: blank moves left, right, up or down.
- goal test: as in figure
- path cost: length of path

Example - The 8-queens problem

The goal of this problem is to place 8 queens on the board so that none can attack the others. The following is not a solution!

- goal test: 8 queens on board, none attacked
- path cost: irrelevant
- states: any arrangement of 0-8 queens on the board
- operators: add or remove a queen to/from any square

Table of Contents

State space graph

A state space graph is a mathematical representation of a search problem.

- Nodes are (abstracted) world configurations
- Arcs represent transitions resulting from actions
- The goal test is a set of goal nodes
- Each state occurs only once

The state space for the vacuum world.

Figure 3.3 The state space for the vacuum world. Links denote actions: L = *Left*, R = *Right*, S = *Suck*.

Table of Contents

Searching for Solutions

- We can use a **search algorithms** to find a solution to a search problem.
- In particular, we can find **a path (i.e. sequence of actions) from the start state to the goal state** in the state space graph.
- Search algorithms can be **uninformed search methods** or **informed search methods.**
- A solution is an **action sequence**, so search algorithms work by considering various possible action sequences.
- The SEARCH TREE possible action sequences starting at the initial state form a search tree with the initial state at the root, the branches are actions and the nodes correspond to states in the state space of the problem.

Searching for Solutions-cont

- By expanding the current state; that is, **applying each legal action to the current state, thereby generating a new set of states**.

- After expanding Arad, add three branches from the parent node, leading to three new child nodes

- After expanding Sibiu these six nodes is a leaf node, The set of all leaf nodes available for expansion at any given point is called the **frontier**. (Many authors call it the open list)

- The process of expanding nodes on the frontier continues until either a solution is found or there are no more states to expand

Searching for Solutions-cont

- The way to avoid exploring redundant paths is to remember, TREE-SEARCH algorithm with a data structure called the explored set (also known as the closed list), which remembers every expanded node(visited).

- Newly generated nodes that match previously generated nodes—ones in the explored set or the frontier—can be discarded instead of being added to the frontier.

- Search algorithms all share this basic structure; they vary primarily according to how they choose which state to expand next—the so-called search strategy.

- The frontier needs to be stored in such a way that the search algorithm can easily choose the next node to expand according to its preferred strategy.

Search strategy and Data Structure

- The appropriate data structure for the search strategy is a queue. Three common variants are:
  - the first-in, first-out or **FIFO queue, which pops the oldest element of the queue;**
  - the last-in, first-out or **LIFO queue (also known as a stack), which pops the newest element of the queue;**
  - the **priority queue**, which pops the element of the queue with the highest priority according to some ordering function.

Partial search trees for finding a route

Figure 3.6 Partial search trees for finding a route from Arad to Bucharest. Nodes that have been expanded are shaded; nodes that have been generated but not yet expanded are outlined in bold; nodes that have not yet been generated are shown in faint dashed lines.

General tree-search and graph-search algorithms

function TREE-SEARCH(*problem*) **returns** a solution, or failure
    initialize the frontier using the initial state of *problem*
    loop do
        if the frontier is empty **then return** failure
        choose a leaf node and remove it from the frontier
        if the node contains a goal state **then return** the corresponding solution
        expand the chosen node, adding the resulting nodes to the frontier

---

function GRAPH-SEARCH(*problem*) **returns** a solution, or failure
    initialize the frontier using the initial state of *problem*
    initialize the explored set to be empty
    loop do
        if the frontier is empty **then return** failure
        choose a leaf node and remove it from the frontier
        if the node contains a goal state **then return** the corresponding solution
        add the node to the explored set
        expand the chosen node, adding the resulting nodes to the frontier
           only if not in the frontier or explored set

Figure 3.7   An informal description of the general tree-search and graph-search algorithms. The parts of GRAPH-SEARCH marked in bold italic are the additions needed to handle repeated states.

Measuring problem-solving performance

- **completeness:** is the strategy guaranteed to find a solution when there is one?

- **optimality:** does the search strategy find the highest quality solution when there are multiple solutions?

- **time complexity:** how long does it take to find a solution?

- **space complexity:** how much memory is required to perform the search?
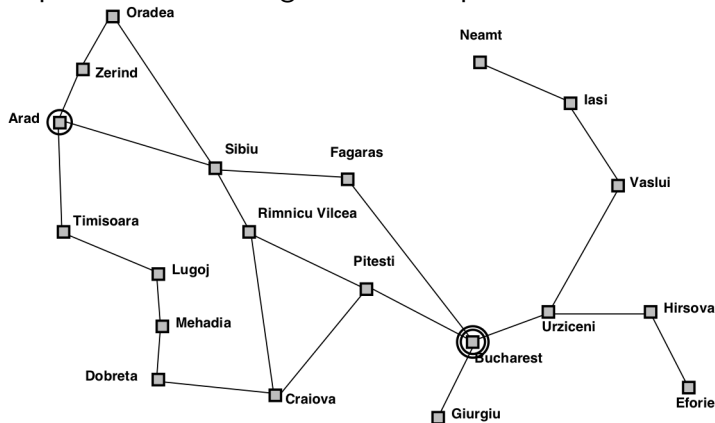
Table of Contents

Breadth-first search

Strategy: expand shallowest node first
Implementation: Fringe is a FIFO queue

Properties of Breadth-first search

Complete: Guaranteed to find a solution if one exists?
Yes (if b is finite)
Optimal: Guaranteed to find the least cost path?
Yes (if cost = 1 per step); not optimal in general
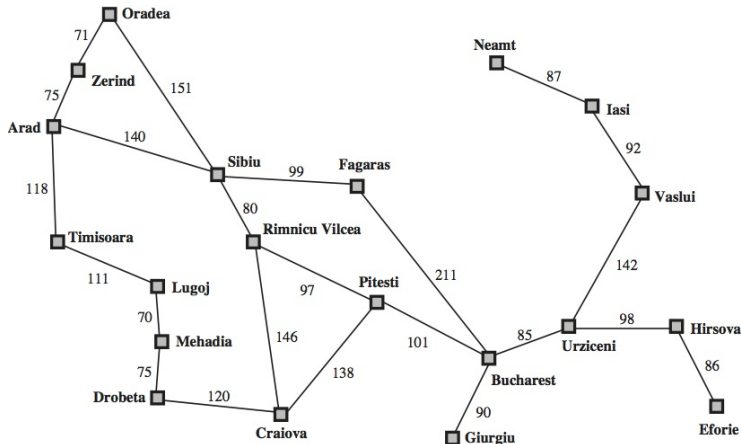**Time:** $1 + b + b^2 + b^3 + \ldots + b^d = O(b^d)$, i.e., exponential in d
Space: $O(b^d)$ (keeps every node in memory)
Space is the big problem; can easily generate nodes at 1MB/sec
so 24hrs = 86GB.

s tiers

1 node

b nodes

b^2 nodes
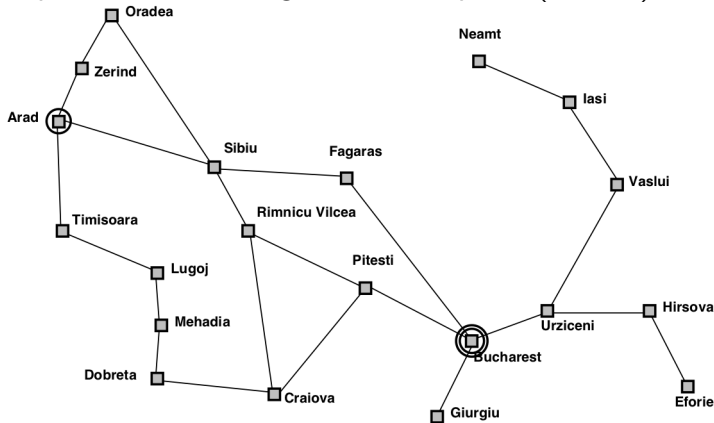
b^s nodes

b^m nodes

Uniform-cost search

- Instead of expanding the shallowest node, uniform-cost search expands the node n with the lowest path cost $g(n)$

Depth-first search

Strategy: expand deepest node first
Implementation: Fringe is a LIFO queue (a stack)

Properties of Depth-first search

Complete: No. fails in infinite-depth spaces, spaces with loops
Modify to avoid repeated states along path
$\Rightarrow$ complete in finite spaces
Optimal: No
Time: $O(b^m)$: terrible if m is much larger than d
but if solutions are dense, may be much faster than breadth-first
Space: $O(bm)$, i.e., linear space!

Depth-limited search

- depth-first search with depth limit L
- **Implementation**: Nodes at depth L have no successors

Iterative deepening search

- Iterative deepening search (or iterative deepening depth-first search) is a general strategy, often used in combination with depth-first tree search, which finds the best depth limit.
- It does this by gradually increasing the limit—first 0, then 1, then 2, and so on—until a goal is found

Bidirectional search

Figure 3.20   A schematic view of a bidirectional search that is about to succeed when a branch from the start node meets a branch from the goal node.

Table of Contents

Comparing uninformed search

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes[a]	Yes[a,b]	No	No	Yes[a]	Yes[a,d]
Time	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes[c]	Yes	No	No	Yes[c]	Yes[c,d]

Figure 3.21   Evaluation of tree-search strategies. b is the branching factor; d is the depth of the shallowest solution; m is the maximum depth of the search tree; l is the depth limit. Superscript caveats are as follows: [a] complete if b is finite; [b] complete if step costs $\geq \epsilon$ for positive ϵ; [c] optimal if step costs are all identical; [d] if both directions use breadth-first search.

Table of Contents

Exercise/Home Work

- Simulate BFS,DFS,IDS,UCS using the graph. Show the expand tree, explored set(visited), Frontier and path cost

Table of Contents

References

📄 Stuart Russell and Peter Norvig. 2009. Artificial Intelligence: A Modern Approach (3rd ed.). Prentice Hall Press, Upper Saddle River, NJ, USA.

📄 https://www.cs.cmu.edu/ 15281/coursenotes/search/

📄 https://inst.eecs.berkeley.edu/ cs188/sp22/